

# Multivariate Analysis

## Classification

Predict the species of a penguin based on their characteristics.

```
In [1]: import pandas as pd
import seaborn as sns
import numpy as np
```

```
In [2]: penguins = sns.load_dataset("penguins")
```


```
In [3]: penguins_no_na = penguins.dropna()
penguins_no_na.head()
```

```
Out[3]:
```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0	Male
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0	Female
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0	Female
4	Adelie	Torgersen	36.7	19.3	193.0	3450.0	Female
5	Adelie	Torgersen	39.3	20.6	190.0	3650.0	Male

```
In [4]: grouped = penguins_no_na.groupby('species')
normalized = (grouped.mean() / grouped.mean().max()).transpose()
errors = (grouped.std() / grouped.mean().max()).transpose()
normalized.plot(kind='bar', yerr=errors)
```

```
Out[4]: <AxesSubplot:>
```



Op basis van bovenstaande grafiek is te verwachten dat species niet goed voorspeld kunnen worden als je maar één numerical gebruikt, er is altijd wel overlap tussen twee species' confidence intervals. Ook is het bij 3 van de 4 dezelfde species die overlap hebben, dus is het nodig om in ieder geval bill\_length\_mm toe te voegen, en ten minste één van de andere 3 numericals.

```
In [5]: from sklearn.model_selection import train_test_split
```

```
In [6]: penguins_train, penguins_test = train_test_split(penguins_no_na, test_size=0.3, random_s
```

```
In [7]: from sklearn.tree import DecisionTreeClassifier
```

```
In [8]: features= ['bill_length_mm', 'flipper_length_mm', 'body_mass_g']
dt_classification = DecisionTreeClassifier(max_depth = 5)
dt_classification.fit(penguins_train[features], penguins_train['species'])
```

```
Out[8]: DecisionTreeClassifier(max_depth=5)
```

Een diepte hoger dan 5 geeft niet meer nodes in de uiteindelijke tree.

```
In [9]: def calculate_accuracy(predictions, actuals):
        if(len(predictions) != len(actuals)):
            raise Exception("The amount of predictions did not equal the amount of actuals")

        return (predictions == actuals).sum() / len(actuals)
```

```
In [10]: predictionsOnTrainset = dt_classification.predict(penguins_train[features])
predictionsOnTestset = dt_classification.predict(penguins_test[features])
```

```
accuracyTrain = calculate_accuracy(predictionsOnTrainset, penguins_train.species)
accuracyTest = calculate_accuracy(predictionsOnTestset, penguins_test.species)
```

```
print("Accuracy on training set " + str(accuracyTrain))
print("Accuracy on test set " + str(accuracyTest))
```

Accuracy on training set 1.0

Accuracy on test set 0.93

De accuracy voor de test set is iets lager, wat niet heel raar is. Dit betekent dat het model niet perfect is, maar over het algemeen wel een vrij goede fit.

```
In [11]: from sklearn import tree
import graphviz

def plot_tree_classification(model, features, class_names):
    # Generate plot data
    dot_data = tree.export_graphviz(model, out_file=None,
                                    feature_names=features,
                                    class_names=class_names,
                                    filled=True, rounded=True,
                                    special_characters=True)

    # Turn into graph using graphviz
    graph = graphviz.Source(dot_data)

    # Write out a pdf
    graph.render("Trees/decision_tree_15")

    # Display in the notebook
    return graph
```

```
In [12]: plot_tree_classification(dt_classification, features, np.sort(penguins.species.unique()))
```

Out[12]:

